

# Speech autoencoder using deep fully-connected networks

Andrew Senior

1/3/19

# Implementation Summary

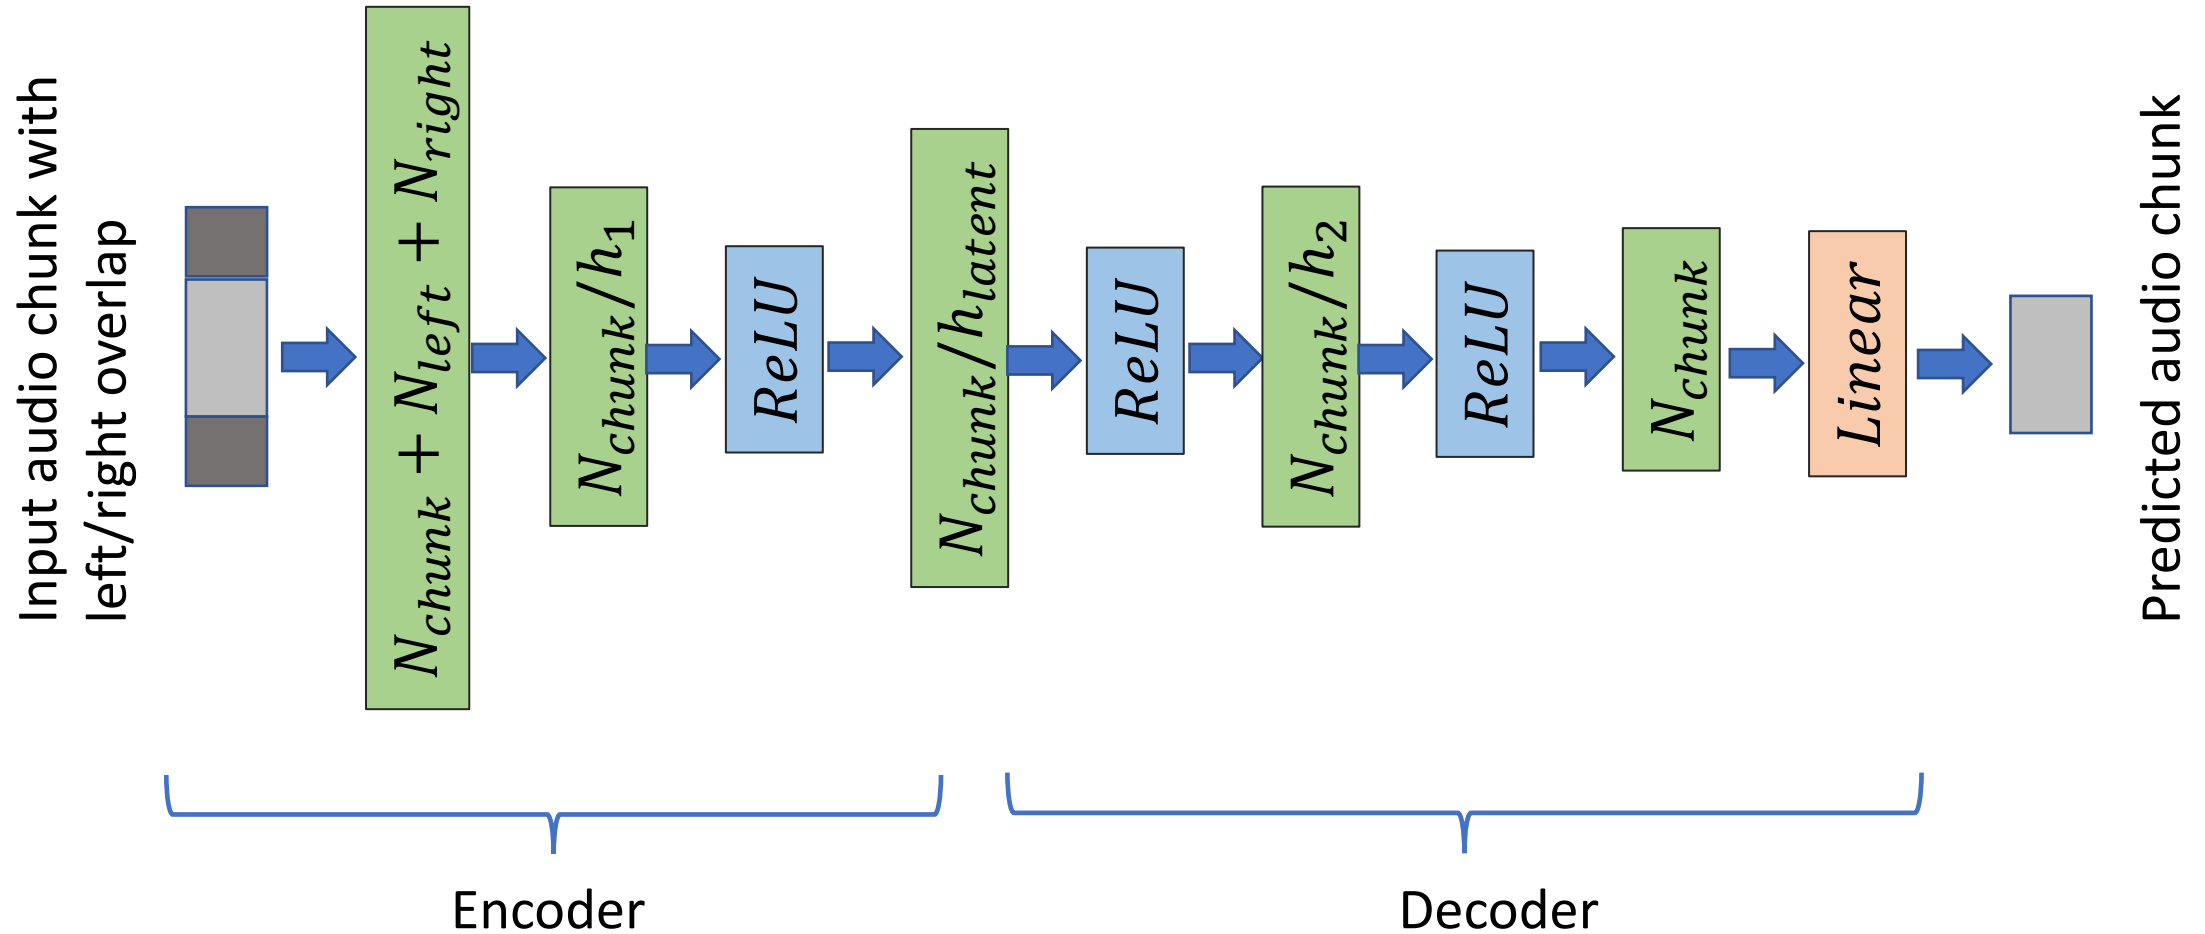
- Vanilla autoencoder built using Keras/TensorFlow
- Custom pipeline for pre-processing and on-the-fly batch generation
- Specialized utility functions for training on LibriSpeech corpuses
- Mini-batches generated by traversing the raw audio signal for each speaker and breaking it into chunks of a specified size which become the feature vectors.
- Overlap allowed on input chunks as a means for incorporating temporal structure into the autoencoder. The corresponding non-overlapping output chunks are used as labels.

# Pre-processing Workflow

## **build\_speech\_dict.py**

- Generates a “speech\_dict” data structure from the corpus
  - Reflects the internal structure of the LibriSpeech corpus. Useful for traversing the speakers, chapters, and utterances.
- Concatenates each speaker’s utterances into a master 1D numpy array for the speaker. Writes array to disk.
  - The sequence of the utterances is preserved in the array.
  - Speaker array used for fast on-the-fly batch generation



# Example Network Architecture



# Training Experiments

- dev-clean dataset (40 speakers, 5.1 hrs total, 294M samples @16kHz)
- Tried chunk size, batch size from Chorowski, et al. (2019)
  - $N_{chunk} = 5120$  ( $\Delta t_{chunk} = 320ms$ , 16kHz), batch size = 64
  - This chunk size is too large for the vanilla autoencoder and training fails with a non-decreasing loss function
- Reducing chunk size and increasing batch size proved successful
  - $N_{chunk} = 800$  ( $\Delta t_{chunk} = 50ms$ , 16kHz), batch size = 128
- Training parameters:
  - 100 epochs, learning rate =  $1e-4$ , Adam optimizer, MSE loss

# 3-layer networks




$$N_{chunk} \rightarrow \frac{N_{chunk}}{h_{latent}} \rightarrow N_{chunk}$$



# 3-layer networks

Speaker 2803

250ms audio sample (5 chunks)



# 5-layer networks




$$N_{chunk} \rightarrow \frac{N_{chunk}}{h_1} \rightarrow \frac{N_{chunk}}{h_{latent}} \rightarrow \frac{N_{chunk}}{h_1} \rightarrow N_{chunk}$$





# 5-layer networks

Speaker 2803

250ms audio sample (5 chunks)



# Effect of network depth




Symmetric architectures,  $h_{latent}=16$ , size reduction only in the latent layer

# Effect of network depth

Speaker 2803

250ms audio sample (5 chunks)



Symmetric architectures,  $h_{latent}=16$ , size reduction only in the latent layer